

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.arima_model import ARIMA
import tensorflow as tf
from statsmodels.tsa.arima.model import ARIMA
import warnings
warnings.filterwarnings("ignore")
from statsmodels.tsa.arima_model import ARIMA
from statsmodels.tsa.arima.model import ARIMA

import matplotlib.pyplot as plt
import pandas as pd
import statsmodels.api as sm
from statsmodels.graphics.tsaplots import plot_predict
from statsmodels.tsa.arima.model import ARIMA

```

```
df = pd.read_csv("/content/drive/MyDrive/TM/a10.csv", index_col = 'date')
```

```
df
```

	value
date	
1991-07-01	3.526591
1991-08-01	3.180891
1991-09-01	3.252221
1991-10-01	3.611003
1991-11-01	3.565869
...	...
2008-02-01	21.654285
2008-03-01	18.264945
2008-04-01	23.107677
2008-05-01	22.912510
2008-06-01	19.431740

204 rows × 1 columns

EDA

```
df.shape
```

```
(204, 1)
```

```
df=df.dropna()
```

```
df.isnull().sum()
```

```
value    0
dtype: int64
```

```
df.head(5)
```

```

            value
date
df.tail()

```

```

            value
date
2008-02-01  21.654285
2008-03-01  18.264945
2008-04-01  23.107677
2008-05-01  22.912510
2008-06-01  19.431740


```

## Visualize of Data

```

df.plot()
# seasonlity data

```



```

# test for stationarity
from statsmodels.tsa.stattools import adfuller

```

```

# check stationary or not
def adfuller_test(value):
    result = adfuller(value)
    labels = ['ADF Test Statistics', 'p-value', 'Number of observation used']
    for value, label in zip(result, labels):
        print(label+': ' +str(value))
    if result[1]<=0.05:
        print("strong evidence against the null hypothesis, reject the null hypothesis")
    else:
        print("weak evidence against the null hypothesis")

```

```

adfuller_test(df['value'])

ADF Test Statistics : 3.14518568930673
p-value : 1.0
Number of observation used : 15
weak evidence against the null hypothesis

```

## Differencing

```

from statsmodels.graphics.tsaplots import plot_acf, plot_pacf

plt.rcParams.update({'figure.figsize':(9,7), 'figure.dpi':120})


# Original Series
fig, axes = plt.subplots(3, 2, sharex=True)
axes[0, 0].plot(df.value); axes[0, 0].set_title('Original Series')
plot_acf(df.value, ax=axes[0, 1])

```

```
# 1st Differencing
axes[1, 0].plot(df.value.diff()); axes[1, 0].set_title('1st Order Differencing')
plot_acf(df.value.diff().dropna(), ax=axes[1, 1])

# 2nd Differencing
axes[2, 0].plot(df.value.diff().diff()); axes[2, 0].set_title('2nd Order Differencing')
plot_acf(df.value.diff().diff().dropna(), ax=axes[2, 1])

plt.show()
```



```
# PACF plot of 1st differenced series
plt.rcParams.update({'figure.figsize':(9,3), 'figure.dpi':120})

fig, axes = plt.subplots(1, 2, sharex=True)
axes[0].plot(df.value.diff()); axes[0].set_title('1st Differencing')
axes[1].set_ylim=(0,5)
plot_pacf(df.value.diff().dropna(), ax=axes[1])

plt.show()
```

```
/usr/local/lib/python3.9/dist-packages/statsmodels/graphics/tsaplots.py:348: FutureWarning: The default method 'yw' can pro
warnings.warn(
```


## 1st Differencing

## Partial Autocorrelation

```
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
plt.rcParams.update({'figure.figsize':(9,3), 'figure.dpi':120})

fig, axes = plt.subplots(1, 2, sharex=True)
axes[0].plot(df.value.diff()); axes[0].set_title('1st Differencing')
axes[1].set(ylim=(0,1.2))
plot_acf(df.value.diff().dropna(), ax=axes[1])

plt.show()
```



```
import statsmodels.api as sm
```

```
import statsmodels.api as smapi
```

```
# 1,1,2 ARIMA Model
model = ARIMA(df.value, order=(1,1,2))
model_fit = model.fit()
print(model_fit.summary())
```

```
SARIMAX Results
=====
Dep. Variable:          value    No. Observations:          204
Model:                 ARIMA(1, 1, 2)    Log Likelihood        -424.570
Date:                 Wed, 15 Mar 2023    AIC                   857.140
Time:                 08:49:23    BIC                   870.393
Sample:              07-01-1991    HQIC                  862.502
                        - 06-01-2008
Covariance Type:      opg
=====
              coef    std err          z      P>|z|      [0.025    0.975]
-----
ar.L1         0.4178     0.356      1.174     0.240     -0.280     1.115
ma.L1        -0.9546     0.377     -2.531     0.011     -1.694    -0.215
ma.L2         0.0969     0.272     0.356     0.722     -0.437     0.631
sigma2        3.8259     0.269     14.209     0.000      3.298     4.354
=====
Ljung-Box (L1) (Q):           0.46    Jarque-Bera (JB):          135.61
Prob(Q):                     0.50    Prob(JB):                 0.00
Heteroskedasticity (H):       9.82    Skew:                     -0.80
Prob(H) (two-sided):          0.00    Kurtosis:                  6.67
=====
```

```
Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-step).
```

```
# 1,1,1 ARIMA Model
model = ARIMA(df.value, order=(1,1,1))
model_fit = model.fit()
print(model_fit.summary())
```

```
SARIMAX Results
=====
Dep. Variable:          value    No. Observations:          204
```

```

Model:          ARIMA(1, 1, 1)    Log Likelihood    -424.762
Date:           Wed, 15 Mar 2023  AIC              855.524
Time:           08:50:00          BIC              865.463
Sample:         07-01-1991        HQIC              859.545
              - 06-01-2008

Covariance Type: opg
=====
              coef    std err          z      P>|z|      [0.025    0.975]
-----
ar.L1          0.3009     0.094       3.195     0.001     0.116     0.485
ma.L1         -0.8300     0.048    -17.204     0.000    -0.925    -0.735
sigma2         3.8327     0.259     14.790     0.000     3.325     4.341
=====
Ljung-Box (L1) (Q):           0.72    Jarque-Bera (JB):           130.26
Prob(Q):                     0.40    Prob(JB):                0.00
Heteroskedasticity (H):       9.98    Skew:                   -0.75
Prob(H) (two-sided):          0.00    Kurtosis:                6.63
=====

```


Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```

# Plot residual errors
residuals = pd.DataFrame(model_fit.resid)
fig, ax = plt.subplots(1,2)
residuals.plot(title="Residuals", ax=ax[0])
residuals.plot(kind='kde', title='Density', ax=ax[1])
plt.show()

```



```

# Actual vs Fitted
model_fit.plot_predict()
plt.show()

```

```

-----
AttributeError                                Traceback (most recent call last)
<ipython-input-77-1705293967f2> in <module>
      1 # Actual vs Fitted
----> 2 model_fit.plot_predict()
      3 plt.show()

/usr/local/lib/python3.9/dist-packages/statsmodels/base/wrapper.py in __getattr__(self, attr)
     32     pass
     33
----> 34     obj = getattr(results, attr)
     35     data = results.model.data
     36     how = self._wrap_attrs.get(attr)

AttributeError: 'ARIMAResults' object has no attribute 'plot_predict'

```

SEARCH STACK OVERFLOW

```

from statsmodels.tsa.stattools import acf

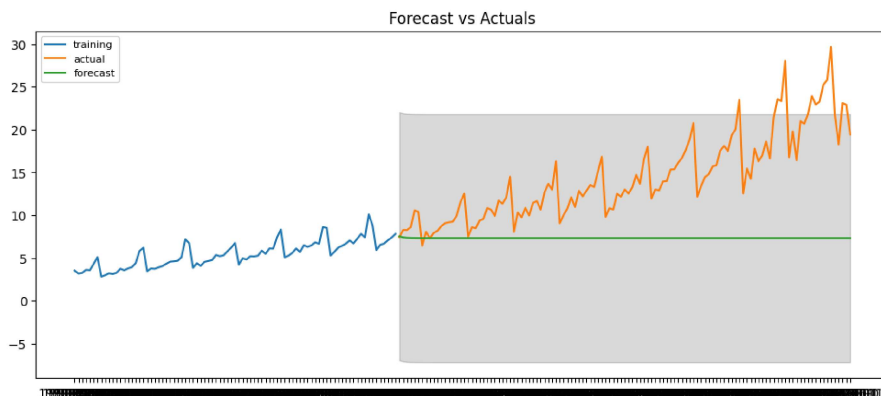
# Create Training and Test
train = df.value[:85]
test = df.value[85:]

```

```
# Build Model
# model = ARIMA(train, order=(3,2,1))
model = ARIMA(train, order=(1, 1, 1))
fitted = model.fit()

# Forecast
fc = fitted.forecast(119, alpha=0.05)
fc_series = pd.Series(fc, index=test.index)
se = fitted.forecast(119, alpha=0.05)[1]
lower_series = pd.Series(fc - 1.96*se, index=test.index)
upper_series = pd.Series(fc + 1.96*se, index=test.index)

# Plot
plt.figure(figsize=(12,5), dpi=100)
plt.plot(train, label='training')
plt.plot(test, label='actual')
plt.plot(fc_series, label='forecast')
plt.fill_between(lower_series.index, lower_series, upper_series,
                 color='k', alpha=.15)
plt.title('Forecast vs Actuals')
plt.legend(loc='upper left', fontsize=8)
plt.show()
```



```
# Build Model
model = ARIMA(train, order=(3, 2, 1))
fitted = model.fit()
print(fitted.summary())

# Forecast
fc = fitted.forecast(119, alpha=0.05)
fc_series = pd.Series(fc, index=test.index)
se = fitted.forecast(119, alpha=0.05)[1]
lower_series = pd.Series(fc - 1.96*se, index=test.index)
upper_series = pd.Series(fc + 1.96*se, index=test.index)

# Plot
plt.figure(figsize=(12,5), dpi=100)
plt.plot(train, label='training')
plt.plot(test, label='actual')
plt.plot(fc_series, label='forecast')
plt.fill_between(lower_series.index, lower_series, upper_series,
                 color='k', alpha=.15)
plt.title('Forecast vs Actuals')
plt.legend(loc='upper left', fontsize=8)
plt.show()
```

## SARIMAX Results

```

=====
Dep. Variable:          value      No. Observations:          85
Model:                 ARIMA(3, 2, 1)  Log Likelihood         -116.886
Date:                 Wed, 15 Mar 2023  AIC                   243.771
Time:                 09:37:25         BIC                   255.865
Sample:              07-01-1991       HQIC                  248.630
                             - 07-01-1998
Covariance Type:      opg
=====

```

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	-0.2205	0.194	-1.139	0.255	-0.600	0.159
ar.L2	-0.3436	0.141	-2.444	0.015	-0.619	-0.068
ar.L3	-0.0937	0.252	-0.371	0.710	-0.588	0.401
ma.L1	-0.9998	14.653	-0.068	0.946	-29.720	27.720
sigma2	0.9141	13.333	0.069	0.945	-25.219	27.047

```


=====
Ljung-Box (L1) (Q):          0.09  Jarque-Bera (JB):          24.64
Prob(Q):                   0.76  Prob(JB):              0.00
Heteroskedasticity (H):     1.53  Skew:                -0.81
Prob(H) (two-sided):       0.27  Kurtosis:             5.12
=====

```

## Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

## Forecast vs Actuals



## # Accuracy metrics

```

def forecast_accuracy(forecast, actual):
    mape = np.mean(np.abs(forecast - actual)/np.abs(actual)) # MAPE
    me = np.mean(forecast - actual) # ME
    mae = np.mean(np.abs(forecast - actual)) # MAE
    mpe = np.mean((forecast - actual)/actual) # MPE
    rmse = np.mean((forecast - actual)**2)**.5 # RMSE
    corr = np.corrcoef(forecast, actual)[0,1] # corr
    mins = np.amin(np.hstack([forecast[:,None],
                              actual[:,None]]), axis=1)
    maxs = np.amax(np.hstack([forecast[:,None],
                              actual[:,None]]), axis=1)
    minmax = 1 - np.mean(mins/maxs) # minmax
    acf1 = acf(fc-test)[1] # ACF1
    return({'mape':mape, 'me':me, 'mae': mae,
            'mpe': mpe, 'rmse':rmse, 'acf1':acf1,
            'corr':corr, 'minmax':minmax})

```

```
forecast_accuracy(fc, test.values)
```

```

{'mape': 0.23266899442329508,
 'me': -3.880721387967784,
 'mae': 3.9817524734401397,
 'mpe': -0.2193087744854135,
 'rmse': 5.307216001549148,
 'acf1': nan,
 'corr': 0.8736884585142838,
 'minmax': 0.23187003957555852}

```

